

Where Will Your Paper Go? Predicting Journal Submission and Citations with Topic Models

Zhe Guan^{a,1}

This manuscript was compiled on January 9, 2025

Academic impact prediction is an important task for scholars, funding agencies, and journal editors to guide funding allocations, research priorities, and editorial decisions. This study presents a framework that uses the paper's submission year, page count, views, and title abstracts and combines topic modelling techniques with various machine learning classifiers while integrating multiple clustering algorithms to predict both the number of citations a paper will receive and the most suitable journal for submission by excluding variables unavailable prior to submission. Additionally, interactive dashboard tools are provided, allowing users to input paper abstracts and obtain predictions using different models.

Topic Modeling | Machine Learning | Citation Prediction | Journal Classification

Introduction

The exponential growth in the number of academic papers in different journals and on various topics brings unignorable challenges for researchers and policymakers who arrange funding and research priorities in institutions(1). There is an increasing need for intelligent tools that can analyse and predict both citation potential and appropriate target journals for academic papers(2).

In this work, we explored the situation by focusing on the data from three journals within the field of operations research and systems analysis from 2020 to 2022: the *Journal of the Operational Research Society*, *Health Systems*, and *Journal of Simulation*, exploring the practicability of predicting citations and target journals by analysing topics from title and abstract inputs along with other submission parameters. The project is divided into five parts: exploratory data analysis, topic modelling, regression, and classification.

1. Dataset and Cleaning

The analysed dataset in this work includes the 4141 records. The initial features are shown in Table 1.

Table 1. The initial features in dataset from journals

Features	Description
Title	Title of the paper
Journal	Name of the journal
Year	Published year
Pages	Number of pages in the paper
Authors	Names of the authors
Views	Number of views the paper received
Citations	Number of times the paper was cited
Altmetric Score	Altmetric score indicating social and public engagement
Abstract	Summary of the paper's content

Since the title name may reveal useful information for topic modelling, the two variables, "Title" and "Abstract," are combined into a single input variable, "Title Abstract". And "Authors" are

transformed into "Author_Counts" which can provide useful numeric information. Different techniques are applied to fill the missing "Views" values in the train and test data, with details provided in the regression section only where we use "View" and "altmetric Score" variables. Considering that one of our targets is predicting journal names, the two variables, the "Views" and "Almetric Score," are excluded from the classification model, as these metrics are not available prior to submission. A unique identifier, "row_id", was added to each record. "Title_Abstract" was transformed into lemmatisation to standardise word forms first. Next, stop words and numeric tokens were removed to reduce noise in the dataset.

2. Topic Modelling

Bag of words and Comparisons.

The combined text of "Title_Abstract" was tokenised into individual words and bigrams using a bag-of-words(BoW) approach (3). The top 100 frequencies of unigrams and bigrams in the dataset are shown in Figure 1. The bigrams can provide more useful and distinctive information compared to common unigrams. Bigrams are also used in all other parts of the work.

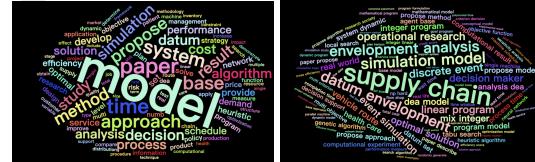


Fig. 1. Top 100 frequencies of unigrams and bigrams from datasets

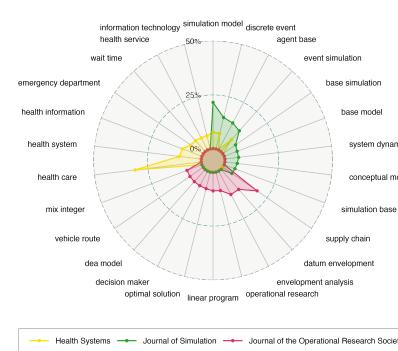


Fig. 2. The top 10 most frequent bigrams across three journals.

Author affiliations: ^aThe University of Southampton

¹ E-mail: zg2u24soton.ac.uk

The BoW frequency of bigrams is analysed from two perspectives: its variation across different journals and over four time periods, as shown in Figure 2 and Figure 3.

As shown in Figure 2, Different areas are focused on across journals, and the high-frequency bigrams show obvious distinctions, which can help in developing future classification models to identify the most appropriate target journal. Some common bigrams still exist across journals, such as “simulation model” and “event simulation” appearing between the journal *Health Systems* and *Journal of Simulation*, as well as “supply chain” shared between *Journal of Simulation* and *Journal of Operational Research Society*.

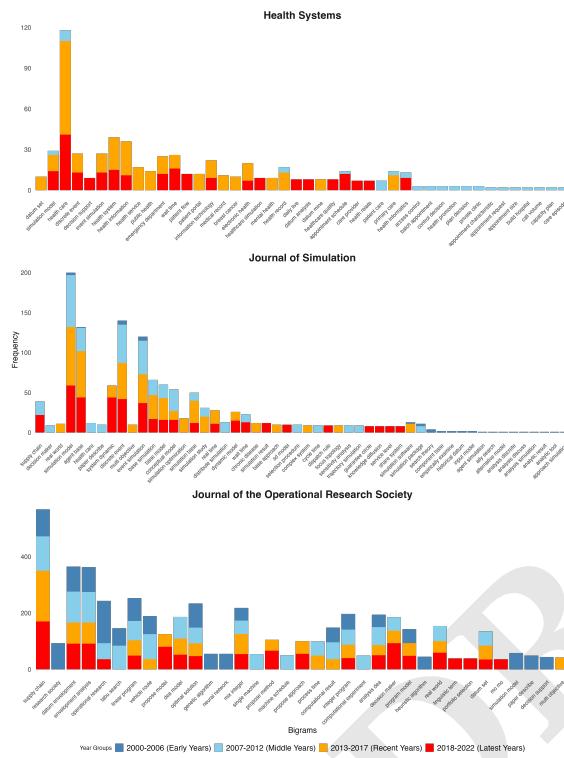


Fig. 3. The top 30 most frequent bigrams from the dataset across four different time periods: 2000–2006, 2007–2012, 2013–2017, and 2018–2022 for three journals.

Some interesting phenomena appear over time across three journals as shown in Figure 3. Firstly, there is no paper published in the journal *Health Systems* before the Middle period (2007–2012), and “health care” is the main topic of the journal and shows increasing popularity across years. Secondly, the key bigrams in the *Journal of Simulation* are relatively concentrated and increasingly popular across years, such as “simulation model”, “discrete event”, and “event simulation”, whereas the *Journal of the Operational Research Society* shows more flat probabilities across bigrams, with more focus on terms like “supply chain.”

Latent Dirichlet Allocation. Latent Dirichlet Allocation (LDA) (4) is a widely used topic modelling algorithm that identifies hidden topics in a collection of documents based on word co-occurrence patterns. It assumes that each document is a mixture of topics, and each topic is a distribution over words. LDA requires specifying a value of k (the number of topics) in the corpus, as in the example shown in Figure 4 with the number of topics $k = 10$. Determining an appropriate value for k is significant, as it probably impacts the quality and interpretability of the topic modelling results. Perplexity

is a common metric used to evaluate the fit of topic models and provides a reference for selecting the optimal k value in LDA.

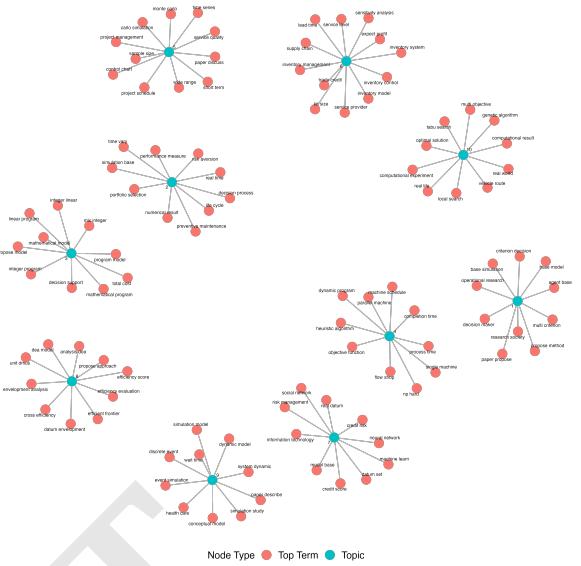


Fig. 4. The identified topics using LDA with the number of topics $k = 10$.

Lower perplexity generally indicates a better model fit. As shown in Figure 5, the 2000–2022 year group has a much higher perplexity compared to other year groups, indicating that over a longer time span, the data is more diverse, and the model finds it harder to capture consistent topics across the entire period. Similarly, the *Journal of the Operational Research Society* has a much higher perplexity compared to other journals, indicating that the journal covers a broader and more diverse range of topics, making it harder for the model to find consistent patterns.

Furthermore, Figure 5 also shows that considering dataset features like year groups can help reduce perplexity and the optimal number value of topics.

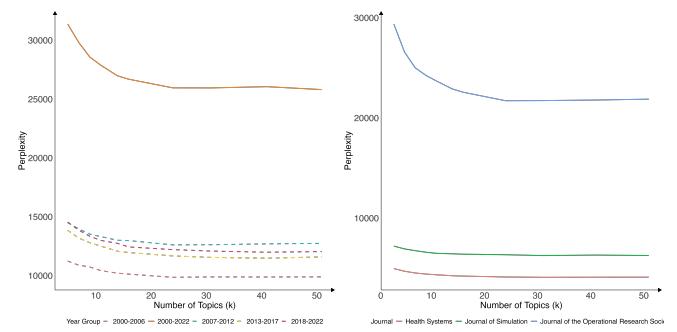


Fig. 5. The perplexity with different numbers of topics across years and journals.

Structural Topic Modelling.

Structural Topic Modelling (STM) (4) provided a method to improve topic modelling performance by incorporating document-level metadata (e.g., publication year, journal name as mentioned above) into the topic modelling process, which builds on the foundation of LDA. Several metrics can be used in STM to evaluate the optimal k (5), such as *Exclusivity*, *Heldout Likelihood*, *Residuals*, and *Semantic Coherence* as shown in Table ??.

Table 2. Comparison of STM Metrics for Different K Values

K	Exclus	Semcoh	Heldout	Residual
2	7.93	-165.02	-8.90	61.93
4	8.82	-184.20	-8.92	47.10
6	9.41	-236.94	-8.97	36.15
8	9.53	-230.01	-8.94	29.66
10	9.67	-219.33	-8.86	24.15
12	9.76	-219.71	-8.84	19.81
14	9.76	-216.37	-8.81	16.59
16	9.80	-223.70	-8.88	14.75
18	9.79	-213.05	-8.84	12.42
20	9.73	-223.25	-8.98	11.30

As shown in Table 2, the STM metrics suggest that $k = 10$ is a reasonable choice, as the majority of the metrics remain stable at this point compared to the stable case of $k = 20$ using LDA shown in Figure 5.

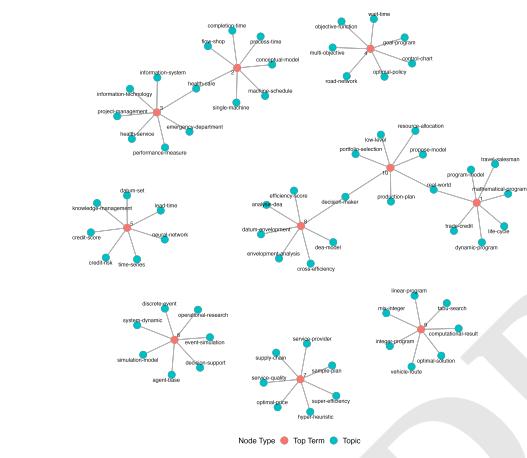


Fig. 6. The identified topics using STM with the number of topics $k = 10$.

One of the benefits of the STM model is that it can easily reveal the topic trend with data features like time. As shown in Figure 7, topics 1, 10, 7, and 8 demonstrate an increasing trend in popularity over the years. In contrast, topics 2, 3, 4, and 5 suggest a potential decline in popularity, although with wide uncertainty bands. Notably, topics 6 and 9 show clear and consistent decreases in popularity over time.

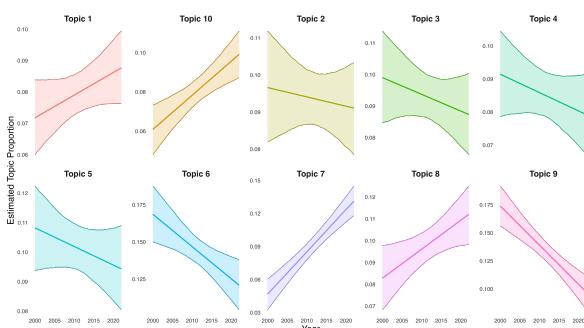


Fig. 7. Topic trends over time with the STM model and topic number $k = 10$.

Although we obtained more information with two different models, the similarity between two models can cause multicollinearity, negatively influencing model training. The similarity S of topics i and j between two models can be evaluated by Formula 1 (6):

$$S_{ij} = \frac{\sum_{w=1}^n A_{wi} B_{wj}}{\sqrt{\sum_{w=1}^n A_{wi}^2} \times \sqrt{\sum_{w=1}^n B_{wj}^2}} \quad [1]$$

Where A_{wi} and B_{wj} are the probabilities $\beta_{wi} = P(\text{word}_w | \text{topic}_i)$ and $\beta_{wj} = P(\text{word}_w | \text{topic}_j)$ of the w^{th} word from the LDA model and STM model.

The similarity matrix shown in Figure 8 suggests that 50% of topics have a similarity greater than 0.7, and 80% of topics have a similarity greater than 0.4 between the two models with the same number of topics $k = 10$. The Principal Component Analysis (PCA) (7) is a common method to obtain uncorrelated principal components and remove correlations.



Fig. 8. The similarity matrix between LDA and STM models with topic number $k = 10$.

Citations Prediction with Regression

Preprocessing.

The topic probabilities generated by the two models were combined with the original data as features for regression analysis. The dataset was split into 80% training data and 20% test data to evaluate model performance. Additionally, the Journal Name column was transformed using one-hot encoding to convert categorical values into binary features.

The missing “Views” values in the training data are filled using predictions from a linear regression model based on ‘Citations,’ while in the test data, they are replaced with mean values.

Linear Regression.

Although two topic models, LDA and STM, have been trained, the similarity may lead to redundancy in the features generated from both models with the same k value. The redundancy can negatively affect regression models (8) due to multicollinearity. PCA is applied to reduce the dimensionality of the combined topic distributions and transform correlated features into a set of uncorrelated principal components. As shown in Figure 9, the first 23 PCA components are selected to train the linear regression model.

XGBoost. eXtreme Gradient Boosting (XGBoost) (9) uses gradient boosting, minimising a loss function by adding new trees that predict the residual errors of the previous trees. However, the effectiveness of this process heavily depends on the selection of appropriate hyperparameters. In this work, a grid search is applied, as shown in Table 3 for optimal choice.

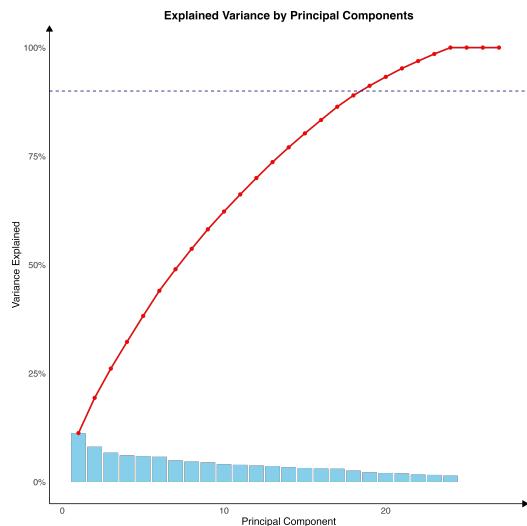


Fig. 9. Explained and cumulative variances by principal components.

RandomForest. Compared to XGBoost, Random Forest (10) takes a different approach by employing an ensemble of independent decision trees. Similar to XGBoost, the performance of Random Forest is sensitive to hyperparameter selection. A grid search is similarly conducted, as shown in Table 3, to identify the best combination of hyperparameters for Random Forest.

Table 3. The grid scan for XGBoost and Random Forest

Model	Parameter	Description	Scan range
XGBoost	nrounds	Number of boosting rounds	100, 300, 500
	η	Learning rate	0.005, 0.01, 0.03, 0.1
	max_depth	Maximum depth of trees	5, 8, 10
	γ	Minimum loss reduction	0, 1
	colsample_bytree	Column subsample ratio	0.8, 1.0
	min_child_weight	Minimum sum of instance weight	1, 5, 10
Random Forest	subsample	Row subsample ratio	0.8, 1.0
	mtry	Number of variables tried at each split	2, 5, 10
	splitrule	Splitting rule	"variance", "extratrees"
	min.node.size	Minimum size of terminal nodes	1, 5, 10

The three models are evaluated using four metrics: mean squared error (MSE), root mean squared error (RMSE), coefficient of determination (R^2), and residuals, as shown in Figure 10.

In all three models, the residuals between predictions and actual values are generally large, particularly considering that most actual values are below 10. The linear regression model shows the longest tails with a skewed centre value, which can be from the non-linear relationships between the target variable and the predictors, suggesting that the models are missing key information, and additional features or context may be required to achieve more accurate predictions.

3. Classification

Clustering for more information.

Clustering (11) helps capture hidden patterns or latent groupings in the data that are not immediately obvious from the original features. These patterns can provide useful additional information for classification models to make better predictions.

In this work, a combined clustering approach is applied. The Hierarchical Density-Based Spatial Clustering of Applications with Noise (HDBSCAN) model (12) is first used to generate initial clusters. The cluster labels obtained from HDBSCAN are then

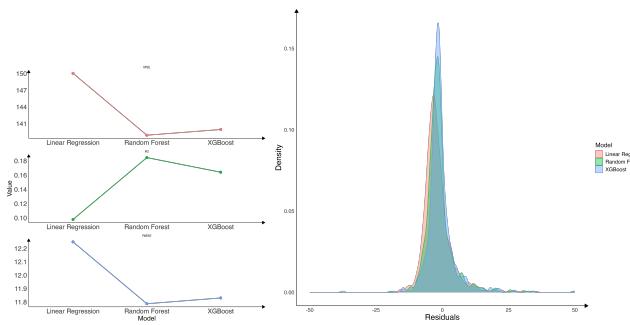


Fig. 10. MSE, RMSE, R^2 , and residual density for regression model evaluations.

used as an input feature for the KMeans method (13) to build a second clustering model. The top three words with the highest mean Term Frequency-Inverse Document Frequency (TF-IDF) (14) scores for each cluster are presented in Table 4. Additionally, the corresponding journal names are displayed in Table 5, which shows that Cluster 3 is associated with the "Journal of the Operational Research Society," while Clusters 1 and 2 provide useful insights for the other two journals.

Table 4. Top Words with Highest TF-IDF Scores by Cluster

Cluster	Word	TF-IDF Score
1	supply chain	0.0105
	operational research	0.0108
	datum envelopment	0.0077
2	supply chain	0.0102
	operational research	0.0090
	simulation model	0.0073
3	research society	0.0290
	operational research	0.0275
	supply chain	0.0092

Table 5. Number of Articles by Journal and Cluster

Cluster	Health Systems Journal	Journal of Simulation	Journal of the Operational Research Society
1	122	164	963
2	86	217	1281
3	0	23	1285

Before classification, the training data is created by selecting 80% of the original dataset, combined with cluster labels, STM and LDA topic predictions, and a binary variable indicating whether the document contains journal-specific frequent bigrams, as shown in Figure 2.

K-Nearest Neighbours. K-Nearest Neighbours (KNN) (15) is a non-parametric, distance-based algorithm commonly used for classification tasks. The class label of a test sample is determined by a majority vote among its 20 nearest neighbours from the training data. The Euclidean distance metric was used to measure the similarity between samples.

Multinomial Logistic Regression. Multinomial Logistic Regression (LG) (16) is a supervised learning algorithm commonly used for multi-class classification, where the model predicts the class label by estimating the probability of each possible class. It extends binary logistic regression to handle scenarios with more than two categories. In this study, predictions were made by selecting the class with the highest predicted probability for each test sample, ensuring that each record was assigned to a single class.

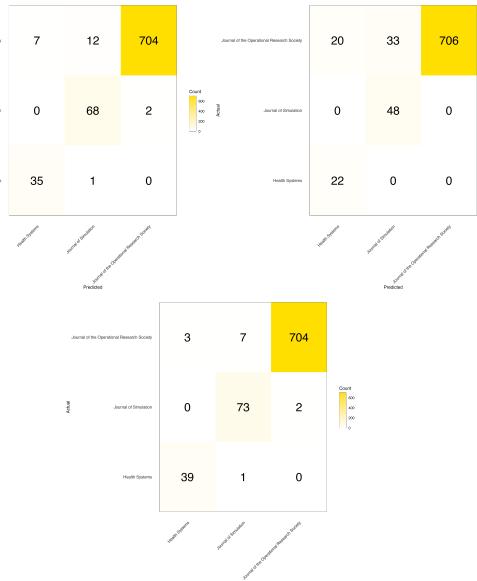


Fig. 11. Confusion matrices for RF, SVM, and XGBoost models.

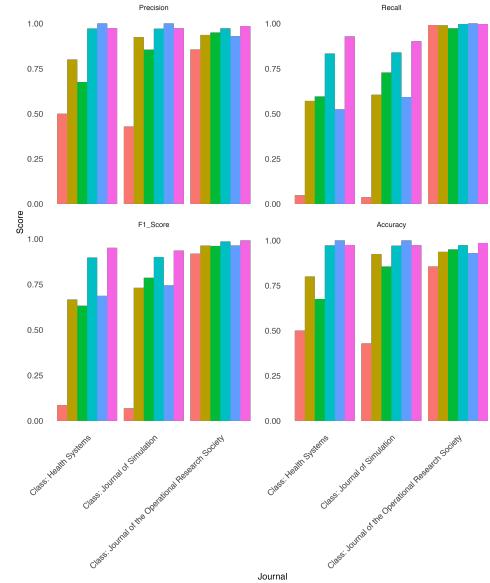


Fig. 12. Metrics for classification evaluation.

Random Forest. The Random Forest model for target journal classification was trained with 500 trees and 3 features randomly selected at each split.

XGBoost. For the classification task, the XGBoost model was trained with a maximum tree depth of 6, a learning rate (η) of 0.3, and 100 boosting rounds. The objective function was set to multi-class classification (softmax) to achieve the multi-class journal prediction.

SVM. Support Vector Machine (SVM) (17) is a supervised learning algorithm that finds an optimal hyperplane to separate different classes in a high-dimensional feature space. In this study, a radial basis function kernel was applied to allow for non-linear decision boundaries, which allows model to capture complex relationships in the data.

Neural Network. A neural network (NN) (18) model is a supervised learning algorithm that models complex relationships between input features and target labels through interconnected layers of neurones. The network architecture in this work consisted of two hidden layers, with 10 neurones in the first layer and 5 neurones in the second layer. The activation function used was sigmoid, ensuring non-linear transformations of the input features.

Evaluations. As shown in Figure 12, in this work, the Random Forest, SVM, and XGBoost models generally show better performance compared to KNN, MLR, and NN algorithms, which have the highest accuracy and precision scores, both exceeding 95%.

Among three models with higher metrics, SVM showed poorer recall and F1 performance, indicating bad prediction for minority samples in the *Health Systems* and *Journal of Simulation* categories compared to the other two models as shown in Figure 11.

4. Dashboard Design

To improve portability, two Shiny-based dashboards are developed for citation count and target journal prediction.

Citation Number Prediction. After receiving input data from users, the previously described cleaning steps and one-hot encoding for journal names are applied. Then, topic probabilities are obtained using LDA and STM, respectively. During the linear regression phase, PCA is first applied to address multicollinearity between the STM and LDA models; additional regression results are generated using trained XGBoost and the Random Forest models, respectively.

Target Journal Prediction. For the target journal prediction, topic modelling results from LDA and STM serve as inputs. Feature engineering is applied to provide a binary variable indicating the presence of frequent journal bigrams. A stacked clustering model combining two clustering algorithms is then applied. Finally, all these features are used as inputs for the XGBoost and Random Forest models in this phase to obtain target fianl prediction shown in the dashboard.

Summary

Providing initial thoughts on which journal is suitable and assessing the impact of a paper is sometimes challenging for scholars and editors. In this work, by comparing three journals focusing on data analytics, healthcare systems, and simulation methodologies from 2000 to 2022, we applied topic modelling and machine learning

621	techniques to predict the target journal and citation count based	683
622	on title and abstract inputs. Two corresponding dashboards were	684
623	developed. The results show that XGBoost and Random Forest	685
624	models with simple learning parameters can effectively identify	686
625		687
626		688
627		689
628	1 MA Hanson, PG Barreiro, P Crosetto, D Brockington, The strain on scientific publishing. <i>Quant. Sci. Stud.</i> 5 , 823–843 (2024).	690
629	2 K Kousha, M Thelwall, Factors associating with or predicting more cited or higher quality journal articles: An annual review of information science and technology (arist) paper. <i>J. Assoc. for Inf. Sci. Technol.</i> 75 , 215–244 (2024).	691
630	3 WA Qader, MM Armeen, BI Ahmed, An overview of bag of words:importance, implementation, applications, and challenges in 2019 International Engineering Conference (IEC). pp. 200–204 (2019).	692
631	4 P Kherwa, P Bansal, Topic modeling: A comprehensive review. <i>EAI Endorsed Transactions on Scalable Inf. Syst.</i> 7 (2019).	693
632	5 ME Roberts, BM Stewart, D Tingley, stm: An R package for structural topic models. <i>J. Stat. Softw.</i> 91 , 1–40 (2019).	694
633	6 J Han, M Kamber, J Pei, 2 - getting to know your data in <i>Data Mining (Third Edition)</i> , The Morgan Kaufmann Series in Data Management Systems, eds. J Han, M Kamber, J Pei. (Morgan Kaufmann, Boston), Third edition edition, pp. 39–82 (2012).	695
634	7 KP F.R.S., Lili. on lines and planes of closest fit to systems of points in space. <i>The London, Edinburgh, Dublin Philos. Mag. J. Sci.</i> 2 , 559–572 (1901).	696
635	8 R Core Team, <i>R: A Language and Environment for Statistical Computing</i> (R Foundation for Statistical Computing, Vienna, Austria), (2024).	697
636	9 T Chen, C Guestrin, Xgboost: A scalable tree boosting system in <i>Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining</i> , KDD '16. (Association for Computing Machinery, New York, NY, USA), p. 785–794 (2016).	698
637	10 L Breiman, Random forests. <i>Mach. Learn.</i> 45 , 5–32 (2001).	699
638	11 M Maechler, P Rousseeuw, A Struyf, M Hubert, K Hornik, <i>cluster: Cluster Analysis Basics and Extensions</i> , (2024) R package version 2.1.8 — For new features, see the 'NEWS' and the 'Changelog' file in the package source).	700
639	12 RJGB Campello, D Moulavi, A Zimek, J Sander, Hierarchical density estimates for data clustering, visualization, and outlier detection. <i>ACM Trans. Knowl. Discov. Data</i> 10 (2015).	701
640	13 X Jin, J Han, <i>K-Means Clustering</i> , eds. C Sammut, GI Webb. (Springer US, Boston, MA), pp. 563–564 (2010).	702
641	14 C Sammut, GI Webb, eds., <i>TF-IDF</i> . (Springer US, Boston, MA), pp. 986–987 (2010).	703
642	15 A Mucherino, PJ Papajorgji, PM Pardalos, <i>k-Nearest Neighbor Classification</i> . (Springer New York, New York, NY), pp. 83–106 (2009).	704
643	16 , <i>Multinomial Logistic Regression</i> . (John Wiley, Sons, Ltd), pp. 109–124 (2016).	705
644	17 M Hearst, S Dumais, E Osuna, J Platt, B Scholkopf, Support vector machines. <i>IEEE Intell. Syst. their Appl.</i> 13 , 18–28 (1998).	706
645	18 WS McCulloch, W Pitts, A logical calculus of the ideas immanent in nervous activity. <i>The bulletin mathematical biophysics</i> 5 , 115–133 (1943).	707
646		708
647		709
648		710
649		711
650		712
651		713
652		714
653		715
654		716
655		717
656		718
657		719
658		720
659		721
660		722
661		723
662		724
663		725
664		726
665		727
666		728
667		729
668		730
669		731
670		732
671		733
672		734
673		735
674		736
675		737
676		738
677		739
678		740
679		741
680		742
681		743
682		744